

```
In [1]: # Exercise Apply k-NN on Amazon reviews dataset

%matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
import scikitplot.metrics as skplt
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

from sklearn.cross_validation import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from sklearn.cross_validation import cross_val_score

from sklearn.decomposition import TruncatedSVD
from sklearn.preprocessing import StandardScaler

from collections import Counter
from sklearn.metrics import accuracy_score
from sklearn import cross_validation
```

```

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

```

```

C:\Users\hemant\Anaconda\lib\site-packages\sklearn\cross_validation.py:
41: DeprecationWarning: This module was deprecated in version 0.18 in f
avor of the model_selection module into which all the refactored classe
s and functions are moved. Also note that the interface of the new CV i
terators are different from that of this module. This module will be re
moved in 0.20.
    "This module will be removed in 0.20.", DeprecationWarning)
C:\Users\hemant\Anaconda\lib\site-packages\gensim\utils.py:1209: UserWa
rning: detected Windows; aliasing chunkize to chunkize_serial
    warnings.warn("detected Windows; aliasing chunkize to chunkize_seria
l")

```

```

In [3]: # using the SQLite Table to read data.
con = sqlite3.connect(r'G:\machine_learning\Real_world_problem_Predict_
rating_given_product_reviews_on_Amazon\amazon\database1.sqlite')

#filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score
!= 3 """, con)

# Give reviews with Score>3 a positive rating, and reviews with a score
<3 a negative rating.

```

```
def partition(x):
    if x < 3:
        return 0
    return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)
```

Number of data points in our data (525814, 10)

Out[3]:

|   | Id | ProductId  | UserId         | ProfileName                              | HelpfulnessNumerator | HelpfulnessDenominator |
|---|----|------------|----------------|--|----------------------|------------------------|
| 0 | 1  | B001E4KFG0 | A3SGXH7AUHU8GW | delmartian                               | 1                    | 1                      |
| 1 | 2  | B00813GRG4 | A1D87F6ZCVE5NK | dll pa                                   | 0                    | 0                      |
| 2 | 3  | B000LQOCH0 | ABXLMWJIXXAIN  | Natalia<br>Corres<br>"Natalia<br>Corres" | 1                    | 1                      |

# Exploratory Data Analysis

## [7.1.2] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [4]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

Out[4]:

|   | Id     | ProductId  | UserId        | ProfileName     | HelpfulnessNumerator | Helpfuln |
|---|--------|------------|---------------|-----------------|----------------------|----------|
| 0 | 78445  | B000HDL1RQ | AR5J8UI46CURR | Geetha Krishnan | 2                    | 2        |
| 1 | 138317 | B000HDOPYC | AR5J8UI46CURR | Geetha Krishnan | 2                    | 2        |

|          | <b>Id</b> | <b>ProductId</b> | <b>UserId</b> | <b>ProfileName</b> | <b>HelpfulnessNumerator</b> | <b>Helpfuln</b> |
|----------|-----------|------------------|---------------|--------------------|-----------------------------|-----------------|
| <b>2</b> | 138277    | B000HDOPYM       | AR5J8UI46CURR | Geetha Krishnan    | 2                           | 2               |
| <b>3</b> | 73791     | B000HDOPZG       | AR5J8UI46CURR | Geetha Krishnan    | 2                           | 2               |
| <b>4</b> | 155049    | B000PAQ75C       | AR5J8UI46CURR | Geetha Krishnan    | 2                           | 2               |

As can be seen above the same user has multiple reviews of the with the same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that

ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8)

ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
In [5]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True,
inplace=False, kind='quicksort', na_position='last')
```

```
In [6]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time",
"Text"}, keep='first', inplace=False)
final.shape
```

Out[6]: (364173, 10)

```
In [7]: #Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

Out[7]: 69.25890143662969

**Observation:-** It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```
In [8]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

display.head()
```

Out[8]:

|   | Id    | ProductId  | UserId         | ProfileName                | HelpfulnessNumerator | Helpfuln |
|---|-------|------------|----------------|----------------------------|----------------------|----------|
| 0 | 64422 | B000MIDROQ | A161DK06JJMCYF | J. E. Stephens<br>"Jeanne" | 3                    | 1        |
| 1 | 44737 | B001EQ55RW | A2V0I904FH7ABY | Ram                        | 3                    | 2        |

```
In [9]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [10]: #Before starting the next phase of preprocessing lets see the number of  
entries left  
print(final.shape)  
  
#How many positive and negative reviews are present in our dataset?  
final['Score'].value_counts()
```

```
(364171, 10)
```

```
Out[10]: 1    307061  
         0     57110  
         Name: Score, dtype: int64
```

### 7.2.3 Text Preprocessing: Stemming, stop-word removal and Lemmatization.

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [11]: # find sentences containing HTML tags
import re
i=0;
for sent in final['Text'].values:
    if (len(re.findall('<.*?>', sent))):
        print(i)
        print(sent)
        break;
    i += 1;
```

6

I set aside at least an hour each day to read to my son (3 y/o). At this point, I consider myself a connoisseur of children's books and this is one of the best. Santa Clause put this under the tree. Since then, we've read it perpetually and he loves it.<br /><br />First, this book taught him the months of the year.<br /><br />Second, it's a pleasure to read. Well suited to 1.5 y/o old to 4+.<br /><br />Very few children's books are worth owning. Most should be borrowed from the library. This book, however, deserves a permanent spot on your shelf. Sendak's best.



```
In [12]: stop = set(stopwords.words('english')) #set of stopwords
sno = nltk.stem.SnowballStemmer('english') #initialising the snowball s
temmer

def cleanhtml(sentence): #function to clean the word of any html-tags
    cleanr = re.compile('<.*?>')
    cleantext = re.sub(cleanr, ' ', sentence)
    return cleantext
def cleanpunc(sentence): #function to clean the word of any punctuation
or special characters
    cleaned = re.sub(r'[?!|\\\'|\"|#]',r'',sentence)
    cleaned = re.sub(r'[\.,|)|(|\\|/]',r' ',cleaned)
    return cleaned
print(stop)
print('*****')
print(sno.stem('tasty'))
```

```
{'it', 'wasn't', 'now', 'until', 'other', 'than', 'couldn't', 'yourself
s', 'you've', 'here', 'below', 'why', 'ain', 'when', 'mightn', 'won',
'should've', 'hasn't', 'herself', 'shan't', 'that', 'few', 'only', 'a',
'can', 'will', 'on', 'being', 'its', 'because', 'to', 'won't', 'after',
'each', 'your', 'you're', 'itself', 'as', 'did', 'in', 'himself', 'dow
n', 'nor', 'off', 'very', 'the', 'his', 'weren', 'mustn', 'needn', 'is
n't', 'this', 'you', 'again', 'we', 'what', 'were', 'some', 'myself',
'haven't', 'my', 's', 'should', 'don', 'these', 'not', 'further', 'by',
'both', 'yours', 'theirs', 'then', 'out', 'where', 'is', 'with', 'ther
e', 'they', 'do', 'hadn't', 'for', 'wouldn', 'd', 'be', 'or', 'same',
'against', 'been', 'shan', 'shouldn't', 'don't', 'shouldn', 'doesn't',
'me', 'wasn', 'ours', 'all', 'more', 'was', 'couldn't', 'weren't', 'the
mselves', 'her', 'and', 'of', 'from', 've', 'before', 'y', 'most', 'hav
en', 'aren't', 'just', 'wouldn't', 'needn't', 'you'd', 'but', 'having',
'didn', 'does', 'o', 'hasn', 'through', 'aren', 'didn't', 'yourself',
'between', 'which', 'ourselves', 'he', 'any', 'about', 'such', 'own',
'it's', 'them', 'have', 'above', 'once', 'too', 'whom', 'isn', 'under',
're', 'no', 't', 'doesn', 'an', 'am', 'll', 'while', 'mustn't', 'i', 'h
adn', 'she's', 'over', 'had', 'how', 'him', 'at', 'm', 'hers', 'their',
'she', 'you'll', 'up', 'during', 'who', 'if', 'so', 'our', 'doing', 'ar
e', 'into', 'ma', 'has', 'those', 'mightn't', 'that'll'}
*****
tasti
```

```

In [13]: #Code for implementing step-by-step the checks mentioned in the pre-processing phase
# this code takes a while to run as it needs to run on 500k sentences.
if not os.path.isfile('final.sqlite'):
    final_string=[]
    all_positive_words=[] # store words from +ve reviews here
    all_negative_words=[] # store words from -ve reviews here.
    for i, sent in enumerate(tqdm(final['Text'].values)):
        filtered_sentence=[]
        #print(sent);
        sent=cleanhtml(sent) # remove HTML tags
        for w in sent.split():
            # we have used cleanpunc(w).split(), one more split function here because consider w="abc.def", cleanpunc(w) will return "abc def"
            # if we dont use .split() function then we will be considering "abc def" as a single word, but if you use .split() function we will get "abc", "def"
            for cleaned_words in cleanpunc(w).split():
                if((cleaned_words.isalpha()) & (len(cleaned_words)>2)):

                    if(cleaned_words.lower() not in stop):
                        s=(sno.stem(cleaned_words.lower())).encode('utf8')

                        filtered_sentence.append(s)
                        if (final['Score'].values)[i] == 1:
                            all_positive_words.append(s) #list of all words used to describe positive reviews
                        if (final['Score'].values)[i] == 0:
                            all_negative_words.append(s) #list of all words used to describe negative reviews
                    str1 = b" ".join(filtered_sentence) #final string of cleaned words

                    #print("*****")
                    *****")
                    final_string.append(str1)

            #####---- storing the data into .sqlite file -----#####
            #####

```

```

final['CleanedText']=final_string #adding a column of CleanedText which displays the data after pre-processing of the review
final['CleanedText']=final['CleanedText'].str.decode("utf-8")
    # store final table into an SQLite table for future.
conn = sqlite3.connect('final.sqlite')
c=conn.cursor()
conn.text_factory = str
final.to_sql('Reviews', conn, schema=None, if_exists='replace', \
            index=True, index_label=None, chunksize=None, dtype=None)
conn.close()

with open('positive_words.pkl', 'wb') as f:
    pickle.dump(all_positive_words, f)
with open('negative_words.pkl', 'wb') as f:
    pickle.dump(all_negative_words, f)

```

```

In [14]: if os.path.isfile('final.sqlite'):
        conn = sqlite3.connect('final.sqlite')
        final = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score !=
        3 """, conn)
        conn.close()
    else:
        print("Please the above cell")

```

## function to find optimal k from knn with kd-tree algorithm

```

In [15]: # kd_tree algo

def find_optimal_k_kdtree(x_tr_dataset,y_train,x_cv_dataset,y_cv):
    from sklearn.metrics import f1_score

    #Setup arrays to store training and test accuracies
    neighbors = np.arange(1,9)

```

```

# train_accuracy = np.empty(len(neighbors))
acc = np.empty(len(neighbors))
error = np.empty(len(neighbors))
for i,k in enumerate(neighbors):

    # instantiate learning model (k = 30)
    knn = KNeighborsClassifier(n_neighbors=k, algorithm='kd_tree' )

    # fitting the model on crossvalidation train
    knn.fit(x_tr_dataset, y_train)

    # predict the response on the crossvalidation train
    pred = knn.predict(x_cv_dataset)

    # evaluate CV accuracy
    acc[i] = f1_score(y_cv, pred, average='macro') * float(100)
    #print('\nCV accuracy for k = %d is %d%%' % (i, acc))
    error[i] = 100-acc[i]

# optimal_k = int(min(error))
# print('\nThe optimal number of neighbors is %d.' % optimal_k)

#Generate plot
d = acc.max()
i = np.where(acc == d)
optimal_k = int(neighbors[i])
print("optimal_k is:-", optimal_k)
plt.title('k-NN Varying number of neighbors')
plt.plot(neighbors, error, label='Testing Accuracy')
#plt.plot(neighbors, train_accuracy, label='Training accuracy')
plt.legend()
plt.xlabel('Number of neighbors')
plt.ylabel('Misclassification error')
plt.grid()
plt.show()

return optimal_k

```

## function to find optimal k from knn with brute algorithm

```
In [16]: # brute algo

def find_optimal_k_brute(x_tr_dataset,y_train,x_cv_dataset,y_cv):

    from sklearn.metrics import f1_score

    #Setup arrays to store training and test accuracies
    neighbors = np.arange(1,9)
    # train_accuracy = np.empty(len(neighbors))
    acc = np.empty(len(neighbors))
    error = np.empty(len(neighbors))
    for i,k in enumerate(neighbors):

        # instantiate learning model (k = 30)
        knn = KNeighborsClassifier(n_neighbors=k,algorithm='brute' )

        # fitting the model on crossvalidation train
        knn.fit(x_tr_dataset, y_train)

        # predict the response on the crossvalidation train
        pred = knn.predict(x_cv_dataset)

        # evaluate CV accuracy
        acc[i] = f1_score(y_cv, pred, average='macro') * float(100)
        #print('\nCV accuracy for k = %d is %d%%' % (i, acc))
        error[i] = 100-acc[i]

    # optimal_k = int(min(error))
    # print('\nThe optimal number of neighbors is %d.' % optimal_k)

    #Generate plot
    d = acc.max()
    i = np.where(acc == d)
    optimal_k = int(neighbors[i])
```

```

    print("optimal_k is:-", optimal_k)
    plt.title('k-NN Varying number of neighbors')
    plt.plot(neighbors, error, label='Testing Accuracy')
    #plt.plot(neighbors, train_accuracy, label='Training accuracy')
    plt.legend()
    plt.xlabel('Number of neighbors')
    plt.ylabel('Misclassification error')
    plt.grid()
    plt.show()

    return optimal_k

```

## randomly generate data and sort in ascending order

```

In [17]: #random_sample = final.sample(n = 6000)
#random_sample.shape
#random_sample = random_sample.sort_values('Time', axis=0, ascending=True, inplace=False, kind='quicksort', na_position='last')

sorted_sample = final.sort_values('Time', axis=0, ascending=True, inplace=False, kind='quicksort', na_position='last')
sample_60000 = sorted_sample.iloc[0:25000]
sample_60000.shape
y = sample_60000['Score']

```

```

In [18]: # sort the data in 60:20:20 ratio
x_train_size = int(len(sample_60000)*.60)
y_size = int(len(y)*.60)

# split into Train and Test sets
x_train = sample_60000[0:x_train_size]
x_test = sample_60000[x_train_size:len(sample_60000)]

print("total data", len(sample_60000))
print("x_train data", len(x_train))

```

```

#print("x_test data",len(x_test))

y_train = y[0:y_size]
y_test = y[y_size:len(y)]

print("total output data",len(y))
print("total y_train data",len(y_train))
#print("total y_test data",len(y_test))

x_tr_size = int(len(x_test)*.50)
y_tr_size = int(len(y_test)*.50)

x_cv = x_test[0:x_tr_size]
x_ts = x_test[x_tr_size:len(x_test)]

#print("total data",x_tr_size)
print("x_cv data",len(x_cv))
print("x_ts data",len(x_ts))

y_cv = y[0:y_tr_size]
y_ts = y[y_tr_size:len(y_test)]

#print("total data",y_tr_size)
print("y_cv data",len(y_cv))
print("y_ts data",len(y_ts))

total data 25000
x_train data 15000
total output data 25000
total y_train data 15000
x_cv data 5000
x_ts data 5000
y_cv data 5000
y_ts data 5000

```

```

In [19]: # data dimension reduction
def truncatesvd_reduction(c):
    from sklearn import decomposition

```

```

truncatesd = decomposition.TruncatedSVD()
truncatesd.n_components = 784
svd_data = truncatesd.fit_transform(c)

percentage_var_explained = truncatesd.explained_variance_ratio_ / t
runcatesd.explained_variance_ratio_.sum()

cum_var_explained = np.cumsum(percentage_var_explained)

# Plot the PCA spectrum
plt.figure(1, figsize=(6, 4))

plt.clf()
plt.plot(cum_var_explained, linewidth=2)
plt.axis('tight')
plt.grid()
plt.xlabel('n_components')
plt.ylabel('Cumulative_explained_variance')
plt.show()

```

## [7.2.2] Bag of Words (BoW)

In [19]:

```

#BoW
count_vect = CountVectorizer() #in scikit-learn
x_tr_final_counts = count_vect.fit_transform(x_train['CleanedText'].values)
x_cv_final_counts = count_vect.transform(x_cv['CleanedText'].values)
x_ts_final_counts = count_vect.transform(x_ts['CleanedText'].values)

print("the type of count vectorizer ", type(x_tr_final_counts))
print("the shape of out text BOW vectorizer ", x_tr_final_counts.get_shape())
print("the number of unique words ", x_tr_final_counts.get_shape()[1])

print("the type of count vectorizer ", type(x_cv_final_counts))
print("the shape of out text BOW vectorizer ", x_cv_final_counts.get_shape())

```



```

pe())
print("the number of unique words ", x_cv_final_counts.get_shape()[1])

print("the type of count vectorizer ", type(x_ts_final_counts))
print("the shape of out text BOW vectorizer ", x_ts_final_counts.get_shape())
print("the number of unique words ", x_ts_final_counts.get_shape()[1])

```

```

the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (15000, 16311)
the number of unique words 16311
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (5000, 16311)
the number of unique words 16311
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (5000, 16311)
the number of unique words 16311

```

In [20]: *# dimensionality reduction for bowT train data*

```


sc = StandardScaler(with_mean=False)
c = sc.fit_transform(x_tr_final_counts)
truncatesvd_reduction(c)

```

```

C:\Users\hemant\Anaconda\lib\site-packages\sklearn\utils\validation.py:
475: DataConversionWarning: Data with input dtype int64 was converted to
float64 by StandardScaler.
warnings.warn(msg, DataConversionWarning)
C:\Users\hemant\Anaconda\lib\site-packages\sklearn\utils\validation.py:
475: DataConversionWarning: Data with input dtype int64 was converted to
float64 by StandardScaler.
warnings.warn(msg, DataConversionWarning)

```




Here 500 components explain 80% of the variation in bowT train data

```
In [21]: svd = TruncatedSVD(n_components=500)
x_tr_dataset = svd.fit_transform(c)
x_tr_dataset.shape
```

```
Out[21]: (15000, 500)
```

```
In [22]: # dimensionality reduction for bowT cross validation data
c = sc.fit_transform(x_cv_final_counts)
truncatesvd_reduction(c)
```

```
C:\Users\hemant\Anaconda\lib\site-packages\sklearn\utils\validation.py:
475: DataConversionWarning: Data with input dtype int64 was converted t
o float64 by StandardScaler.
warnings.warn(msg, DataConversionWarning)
C:\Users\hemant\Anaconda\lib\site-packages\sklearn\utils\validation.py:
475: DataConversionWarning: Data with input dtype int64 was converted t
o float64 by StandardScaler.
warnings.warn(msg, DataConversionWarning)
```




Here 500 components explain 80% of the variation in bowT cross validated data

```
In [23]: svd = TruncatedSVD(n_components=500)
x_cv_dataset = svd.fit_transform(c)
x_cv_dataset.shape
```

```
Out[23]: (5000, 500)
```

```
In [24]: # dimensionality reduction for bowT test data
c = sc.fit_transform(x_ts_final_counts)
truncatesvd_reduction(c)
```

```
C:\Users\hemant\Anaconda\lib\site-packages\sklearn\utils\validation.py:
475: DataConversionWarning: Data with input dtype int64 was converted t
o float64 by StandardScaler.
warnings.warn(msg, DataConversionWarning)
C:\Users\hemant\Anaconda\lib\site-packages\sklearn\utils\validation.py:
475: DataConversionWarning: Data with input dtype int64 was converted t
o float64 by StandardScaler.
warnings.warn(msg, DataConversionWarning)
```



Here 500 components explain 80% of the variation in bowT test data

```
In [25]: svd = TruncatedSVD(n_components=500)
x_ts_dataset = svd.fit_transform(c)
x_ts_dataset.shape
```

```
Out[25]: (5000, 500)
```

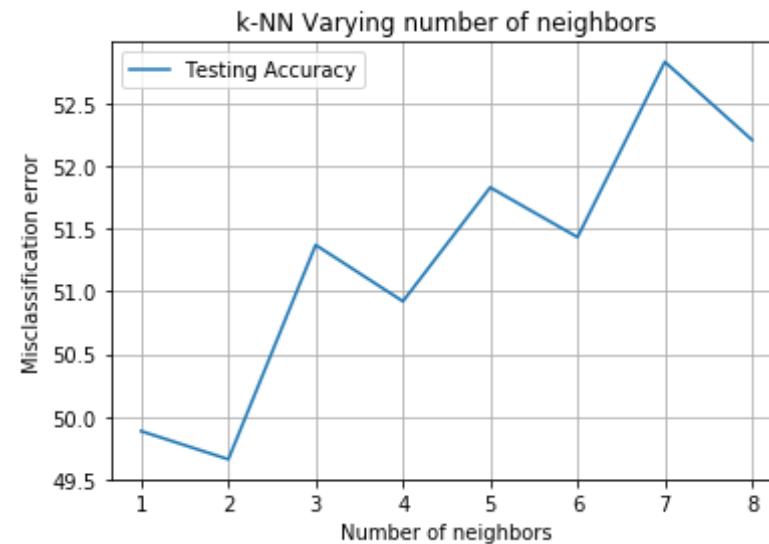
```
In [26]: print(x_tr_dataset.shape)
print(x_cv_dataset.shape)
print(x_ts_dataset.shape)
print(y_train.shape)
print(y_cv.shape)
print(y_ts.shape)
```

```
(15000, 500)
(5000, 500)
(5000, 500)
(15000,)
(5000,)
(5000,)
```

```
In [27]: # Finding Optimal K by simple Cross validation by kd- tree algo

optimal_k = find_optimal_k_kdtree(x_tr_dataset,y_train,x_cv_dataset,y_c
v)
```

optimal\_k is:- 2



```
In [28]: from sklearn.metrics import f1_score
# ===== KNN with k = optimal_k =====
# instantiate learning model k = optimal_k

#knn_optimal = KNeighborsClassifier(n_neighbors = optimal_k)
knn_optimal = KNeighborsClassifier(n_neighbors = 2)

# fitting the model
knn_optimal.fit(x_tr_dataset, y_train)


# predict the response
pred = knn_optimal.predict(x_ts_dataset)
```

```
# evaluate f1_score
f1_bowt_kdtree_optimal_k = optimal_k
f1_bowt_kdtree = f1_score(y_ts, pred, average='macro') * float(100)
print('\nThe f1_score of the knn classifier of KDtree algo for k = %d is %f%%' % (optimal_k, f1_bowt_kdtree))
```

The f1\_score of the knn classifier of KDtree algo for k = 2 is 50.761934%

In [29]: `skplt.plot_confusion_matrix(y_ts, pred)`

Out[29]: `<matplotlib.axes._subplots.AxesSubplot at 0x36c0178128>`



In [30]: `#classification report`  
`from sklearn.metrics import classification_report`  
`print(classification_report(y_ts, pred))`


|          | precision | recall | f1-score | support |
|----------|-----------|--------|----------|---------|
| negative | 0.13      | 0.18   | 0.15     | 554     |
| positive | 0.89      | 0.84   | 0.87     | 4446    |

avg / total      0.81      0.77      0.79      5000

```
In [31]: # Finding Optimal K by simple Cross validation by brute algo

optimal_k = find_optimal_k_brute(x_tr_dataset ,y_train,x_cv_dataset,y_c
v)
```

optimal\_k is:- 2



```
In [32]: from sklearn.metrics import f1_score
# ===== KNN with k = optimal_k =====
# instantiate learning model k = optimal_k

knn_optimal = KNeighborsClassifier(n_neighbors = optimal_k)

# fitting the model
knn_optimal.fit(x_tr_dataset, y_train)

# predict the response
pred = knn_optimal.predict(x_ts_dataset)
```


```
# evaluate f1_score
f1_bowt_brute_optimal_k = optimal_k

f1_bowt_brute = f1_score(y_ts, pred, average='macro') * float(100)
print('\nThe accuracy of the knn classifier of brute algo for k = %d is
%f%%' % (optimal_k, f1_bowt_brute))
```

The accuracy of the knn classifier of brute algo for k = 2 is 50.761934%

In [33]: `skplt.plot_confusion_matrix(y_ts ,pred)`

Out[33]: `<matplotlib.axes._subplots.AxesSubplot at 0x36b7d98c18>`



In [34]: `#classification report`  
`print(classification_report(y_ts, pred))`

|          | precision | recall | f1-score | support |
|----------|-----------|--------|----------|---------|
| negative | 0.13      | 0.18   | 0.15     | 554     |
| positive | 0.89      | 0.84   | 0.87     | 4446    |



avg / total            0.81            0.77            0.79            5000

## [7.2.5] TF-IDF

```
In [20]: tf_idf_vect = TfidfVectorizer(ngram_range=(1,2))
x_tr_final_counts = tf_idf_vect.fit_transform(x_train['CleanedText'].values)
x_cv_final_counts = tf_idf_vect.fit_transform(x_cv['CleanedText'].values)
x_ts_final_counts = tf_idf_vect.fit_transform(x_ts['CleanedText'].values)

print("the type of count vectorizer ",type(x_tr_final_counts))
print("the shape of out text TFIDF vectorizer ",x_tr_final_counts.get_shape())
print("the number of unique words including both unigrams and bigrams "
, x_tr_final_counts.get_shape()[1])


print("the type of count vectorizer ",type(x_cv_final_counts))
print("the shape of out text TFIDF vectorizer ",x_cv_final_counts.get_shape())
print("the number of unique words including both unigrams and bigrams "
, x_cv_final_counts.get_shape()[1])

print("the type of count vectorizer ",type(x_ts_final_counts))
print("the shape of out text TFIDF vectorizer ",x_ts_final_counts.get_shape())
print("the number of unique words including both unigrams and bigrams "
, x_ts_final_counts.get_shape()[1])

the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (15000, 299891)
the number of unique words including both unigrams and bigrams 299891
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (5000, 116881)
```

the number of unique words including both unigrams and bigrams 116881  
the type of count vectorizer <class 'scipy.sparse.csr.csr\_matrix'>  
the shape of out text TFIDF vectorizer (5000, 113417)  
the number of unique words including both unigrams and bigrams 113417

```
In [22]: # dimensionality reduction for TF-IDF train data
sc = StandardScaler(with_mean=False)
c = sc.fit_transform(x_tr_final_counts)
truncatesvd_reduction(c)
```




Here 500 components explain 80% of the variation in TF-IDF train data

```
In [23]: svd = TruncatedSVD(n_components=500)
x_tr_dataset = svd.fit_transform(c)
x_tr_dataset.shape
```

```
Out[23]: (15000, 500)
```

```
In [24]: # dimensionality reduction for TF-IDF cross validate data
c = sc.fit_transform(x_cv_final_counts)
truncatesvd_reduction(c)
```




Here 500 components explain 80% of the variation in TF-IDF cross validate data

```
In [25]: svd = TruncatedSVD(n_components=500)
x_cv_dataset = svd.fit_transform(c)
x_cv_dataset.shape
```

```
Out[25]: (5000, 500)
```

```
In [26]: # dimensionality reduction for TF-IDF test data
c = sc.fit_transform(x_ts_final_counts)
truncatesvd_reduction(c)
```



Here 500 components explain 80% of the variation in TF-IDF test data

```
In [27]: svd = TruncatedSVD(n_components=500)
x_ts_dataset = svd.fit_transform(c)
x_ts_dataset.shape
```

```
Out[27]: (5000, 500)
```

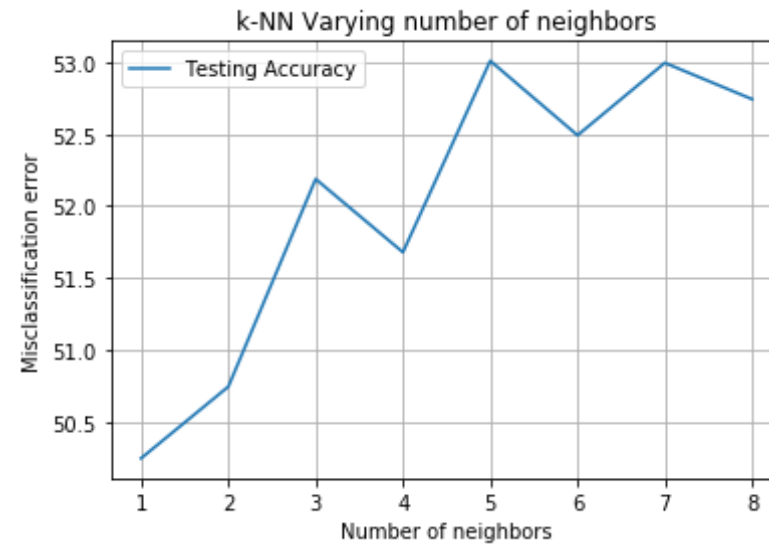
```
In [28]: print(x_tr_dataset.shape)
print(x_cv_dataset.shape)
print(x_ts_dataset.shape)
print(y_train.shape)
print(y_cv.shape)
print(y_ts.shape)
```

```
(15000, 500)
(5000, 500)
(5000, 500)
(15000,)
(5000,)
(5000,)
```

```
In [29]: # Finding Optimal K by simple Cross validation by kd- tree algo

optimal_k = find_optimal_k_kdtree(x_tr_dataset,y_train,x_cv_dataset,y_c
v)
```

optimal\_k is:- 1



```
In [30]: from sklearn.metrics import f1_score
# ===== KNN with k = optimal_k =====
# instantiate learning model k = optimal_k

knn_optimal = KNeighborsClassifier(n_neighbors = optimal_k)

# fitting the model
knn_optimal.fit(x_tr_dataset, y_train)

# predict the response
pred = knn_optimal.predict(x_ts_dataset)


# evaluate f1_score
```

```
f1_tfidf_kdtree_optimal_k = optimal_k
f1_tfidf_kdtree = f1_score(y_ts, pred, average='macro') * float(100)
print('\nThe f1_score of the knn classifier of KDtree algo for k = %d is %f%%' % (optimal_k, f1_tfidf_kdtree))
```

The f1\_score of the knn classifier of KDtree algo for k = 1 is 51.053416%

In [31]: `skplt.plot_confusion_matrix(y_ts, pred)`

Out[31]: `<matplotlib.axes._subplots.AxesSubplot at 0x629cdd9828>`



In [33]: `#classification report`  
`from sklearn.metrics import classification_report`  
`print(classification_report(y_ts, pred))`


|          | precision | recall | f1-score | support |
|----------|-----------|--------|----------|---------|
| negative | 0.13      | 0.13   | 0.13     | 554     |
| positive | 0.89      | 0.89   | 0.89     | 4446    |

avg / total      0.81      0.81      0.81      5000

```
In [34]: # Finding Optimal K by simple Cross validation by brute algo

optimal_k = find_optimal_k_brute(x_tr_dataset ,y_train,x_cv_dataset,y_c
v)
```

optimal\_k is:- 1



```
In [35]: from sklearn.metrics import f1_score
# ===== KNN with k = optimal_k =====
# instantiate learning model k = optimal_k

knn_optimal = KNeighborsClassifier(n_neighbors = optimal_k)

# fitting the model
knn_optimal.fit(x_tr_dataset, y_train)


# predict the response
pred = knn_optimal.predict(x_ts_dataset)
```

```
# evaluate f1_score
f1_tfidf_brute = f1_score(y_ts, pred, average='macro') * float(100)
f1_tfidf_brute_optimal_k = optimal_k
print('\n\nThe f1_score of the knn classifier of brute algo for k = %d is
%f%%' % (optimal_k, f1_tfidf_brute))
```

The accuracy of the knn classifier of brute algo for k = 1 is 51.053416%

In [36]: `skplt.plot_confusion_matrix(y_ts ,pred)`

Out[36]: `<matplotlib.axes._subplots.AxesSubplot at 0x62a2f7b048>`



In [37]: `#classification report`  
`print(classification_report(y_ts, pred))`

|          | precision | recall | f1-score | support |
|----------|-----------|--------|----------|---------|
| negative | 0.13      | 0.13   | 0.13     | 554     |
| positive | 0.89      | 0.89   | 0.89     | 4446    |



avg / total            0.81            0.81            0.81            5000

## [7.2.6] Word2Vec train data

```
In [38]: # Word2Vec for train data
# Train your own Word2Vec model using your own text corpus
i=0
list_of_sent=[]
for sent in x_train['CleanedText'].values:
    list_of_sent.append(sent.split())
```

```
In [39]: print(x_train['CleanedText'].values[0])
print("*****")
print(list_of_sent[0])
```

```
witti littl book make son laugh loud recit car drive along alway sing r
efrain hes learn whale india droop love new word book introduc silli cl
assic book will bet son still abl recit memori colleg
*****
['witti', 'littl', 'book', 'make', 'son', 'laugh', 'loud', 'recit', 'ca
r', 'drive', 'along', 'alway', 'sing', 'refrain', 'hes', 'learn', 'whal
e', 'india', 'droop', 'love', 'new', 'word', 'book', 'introduc', 'sill
i', 'classic', 'book', 'will', 'bet', 'son', 'still', 'abl', 'recit',
'memori', 'colleg']
```

```
In [40]: # min_count = 5 considers only words that occurred atleast 5 times
w2v_model=Word2Vec(list_of_sent,min_count=5,size=50, workers=4)
```

```
In [41]: w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 5 times ",len(w2v_words))
print("sample words ", w2v_words[0:50])
```

```
number of words that occurred minimum 5 times 5603
sample words ['littl', 'book', 'make', 'son', 'laugh', 'loud', 'car',
```

```
'drive', 'along', 'alway', 'sing', 'hes', 'learn', 'india', 'love', 'ne  
w', 'word', 'introduc', 'silli', 'classic', 'will', 'bet', 'still', 'ab  
l', 'memori', 'colleg', 'rememb', 'see', 'show', 'air', 'televis', 'yea  
r', 'ago', 'child', 'sister', 'later', 'bought', 'day', 'thirti', 'some  
th', 'use', 'seri', 'song', 'student', 'teach', 'turn', 'whole', 'schoo  
l', 'purchas', 'children']
```

```
In [42]: w2v_model.wv.most_similar('tasti')
```

```
Out[42]: [('crunchi', 0.8753650188446045),  
( 'satisfi', 0.8572958707809448),  
( 'chewi', 0.8385087251663208),  
( 'nutriti', 0.8355476260185242),  
( 'crispi', 0.834044337272644),  
( 'crisp', 0.8306637406349182),  
( 'yummi', 0.8147004842758179),  
( 'bombay', 0.8044424057006836),  
( 'soft', 0.7914944291114807),  
( 'crunch', 0.7907328605651855)]
```

```
In [43]: w2v_model.wv.most_similar('like')
```

```
Out[43]: [('real', 0.8213435411453247),  
( 'normal', 0.8042137622833252),  
( 'prefer', 0.7925163507461548),  
( 'isnt', 0.7837004661560059),  
( 'think', 0.7791032791137695),  
( 'bland', 0.7584354877471924),  
( 'sort', 0.7490204572677612),  
( 'weird', 0.7471733689308167),  
( 'arent', 0.7392364740371704),  
( 'horribl', 0.7285444140434265)]
```

```
In [44]: # average Word2Vec  
# compute average word2vec for each review.  
sent_vectors = []; # the avg-w2v for each sentence/review is stored in  
this list  
for sent in tqdm(list_of_sent): # for each review/sentence  
    sent_vec = np.zeros(50) # as word vectors are of zero length
```

```

    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
        if cnt_words != 0:
            sent_vec /= cnt_words
    sent_vectors.append(sent_vec)
print(len(sent_vectors))
print(len(sent_vectors[0]))

```

```

100%|████████████████████████████████████████| 15000/15000 [00:25<00:00, 58
0.79it/s]

```

```

15000
50

```

## Word2Vec for cv data

```

In [45]: # Word2Vec
# Train your own Word2Vec model using your own text corpus
i=0
list_of_sent_cv=[]
for sent in x_cv['CleanedText'].values:
    list_of_sent_cv.append(sent.split())

```

```

In [46]: print(x_cv['CleanedText'].values[0])
print("*****")
print(list_of_sent_cv[0])

```

```

realli good hot beverage tran fat havent tri ice would assum would also
good way
*****
['realli', 'good', 'hot', 'beverage', 'tran', 'fat', 'havent', 'tri', 'i
ce', 'would', 'assum', 'would', 'also', 'good', 'way']

```

```
In [47]: # min_count = 5 considers only words that occurred at least 5 times
w2v_model=Word2Vec(list_of_sent_cv,min_count=5,size=50, workers=4)
```

```
In [48]: w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 5 times ",len(w2v_words))
print("sample words ", w2v_words[0:50])
```

```
number of words that occurred minimum 5 times 2984
sample words ['realli', 'good', 'hot', 'beverag', 'tran', 'fat', 'have
nt', 'tri', 'ice', 'would', 'assum', 'also', 'way', 'pretzel', 'think',
'box', 'could', 'pack', 'better', 'deliv', 'care', 'mani', 'broken', 'h
alf', 'small', 'piec', 'two', 'plastic', 'partial', 'crush', 'bought',
'grandson', 'celiac', 'eat', 'absolut', 'delici', 'fool', 'friend', 'fa
mili', 'tast', 'much', 'like', 'regular', 'wheat', 'pasta', 'store', 'w
onder', 'great', 'product', 'recal']
```

```
In [49]: w2v_model.wv.most_similar('tasti')
```

```
Out[49]: [('potato', 0.9988864660263062),
('crunchi', 0.9986565113067627),
('crispi', 0.9983552098274231),
('altern', 0.9982517957687378),
('nut', 0.9981160163879395),
('meal', 0.9980118274688721),
('chip', 0.9978975057601929),
('chewi', 0.9978684782981873),
('veggi', 0.9977338314056396),
('substitut', 0.9977049827575684)]
```

```
In [50]: w2v_model.wv.most_similar('like')
```

```
Out[50]: [('tast', 0.9885746240615845),
('sweet', 0.9875316023826599),
('strong', 0.9857115149497986),
('bitter', 0.9854695796966553),
('realli', 0.983518660068512),
('chocol', 0.9810254573822021),
('light', 0.9804829359054565),
```

```
('hot', 0.9801589846611023),  
('flavor', 0.9797714948654175),  
('make', 0.9789963960647583)]
```

```
In [51]: sent_vectors = []; # the avg-w2v for each sentence/review is stored in  
         # this list  
         for sent in tqdm(list_of_sent_cv): # for each review/sentence  
             sent_vec = np.zeros(50) # as word vectors are of zero length  
             cnt_words = 0; # num of words with a valid vector in the sentence/re  
             view  
             for word in sent: # for each word in a review/sentence  
                 if word in w2v_words:  
                     vec = w2v_model.wv[word]  
                     sent_vec += vec  
                     cnt_words += 1  
             if cnt_words != 0:  
                 sent_vec /= cnt_words  
             sent_vectors.append(sent_vec)  
         print(len(sent_vectors))  
         print(len(sent_vectors[0]))
```

```
100%|████████████████████████████████████████| 5000/5000 [00:05<00:00, 95  
9.87it/s]
```

```
5000  
50
```

## Word2Vec for test data

```
In [52]: # Word2Vec  
         # Train your own Word2Vec model using your own text corpus  
         i=0  
         list_of_sent_ts=[]  
         for sent in x_ts['CleanedText'].values:  
             list_of_sent_ts.append(sent.split())
```

```
In [53]: print(x_ts['CleanedText'].values[0])
```

```
print("*****  
*")  
print(list_of_sent_ts[0])
```

```
german chocol tickl buy everi xmas smooth rich eleg arriv day thank  
*****  
['german', 'chocol', 'tickl', 'buy', 'everi', 'xmas', 'smooth', 'rich',  
'eleg', 'arriv', 'day', 'thank']
```

```
In [54]: # min_count = 5 considers only words that occurred at least 5 times  
w2v_model=Word2Vec(list_of_sent_ts,min_count=5,size=50, workers=4)
```

```
In [55]: w2v_words = list(w2v_model.wv.vocab)  
print("number of words that occurred minimum 5 times ",len(w2v_words))  
print("sample words ", w2v_words[0:50])
```

```
number of words that occurred minimum 5 times 2905  
sample words ['german', 'chocol', 'buy', 'everi', 'xmas', 'smooth', 'r  
ich', 'eleg', 'arriv', 'day', 'thank', 'salt', 'separ', 'rice', 'somet  
h', 'weird', 'happen', 'product', 'sit', 'top', 'tast', 'wick', 'terrib  
l', 'either', 'folk', 'enjoy', 'dont', 'know', 'sushi', 'suppos', 'lik  
e', 'shipment', 'got', 'ruin', 'wont', 'order', 'ever', 'discov', 'sau  
c', 'restaur', 'search', 'everywher', 'garlic', 'avail', 'final', 'foun  
d', 'amazon', 'spici', 'doesnt', 'burn']
```

```
In [56]: w2v_model.wv.most_similar('tasti')
```

```
Out[56]: [('crunch', 0.9992246627807617),  
( 'creami', 0.9990120530128479),  
( 'crunchi', 0.998974084854126),  
( 'salti', 0.9989132285118103),  
( 'salad', 0.9989044666290283),  
( 'aftertast', 0.998854398727417),  
( 'substitut', 0.9988285303115845),  
( 'mapl', 0.998788595199585),  
( 'toast', 0.9987733960151672),  
( 'nut', 0.998724102973938)]
```

```
In [57]: w2v_model.wv.most_similar('like')
```

```
Out[57]: [('tast', 0.9932795763015747),
          ('sweet', 0.9917718172073364),
          ('flavor', 0.9894028902053833),
          ('chocol', 0.9876375198364258),
          ('strong', 0.9838699102401733),
          ('hot', 0.9814661145210266),
          ('milk', 0.9799574017524719),
          ('bitter', 0.9756320714950562),
          ('dark', 0.9754601716995239),
          ('drink', 0.9754283428192139)]
```

```
In [58]: sent_vectors = []; # the avg-w2v for each sentence/review is stored in
         # this list
         for sent in tqdm(list_of_sent_ts): # for each review/sentence
             sent_vec = np.zeros(50) # as word vectors are of zero length
             cnt_words = 0; # num of words with a valid vector in the sentence/review
             for word in sent: # for each word in a review/sentence
                 if word in w2v_words:
                     vec = w2v_model.wv[word]
                     sent_vec += vec
                     cnt_words += 1
             if cnt_words != 0:
                 sent_vec /= cnt_words
             sent_vectors.append(sent_vec)
         print(len(sent_vectors))
         print(len(sent_vectors[0]))
```

```
100%|████████████████████████████████████████| 5000/5000 [00:04<00:00, 101
3.99it/s]
```

```
5000
50
```

```
In [59]: # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
         model = TfidfVectorizer()
         x_tr_final_counts = model.fit_transform(x_train['CleanedText'].values)
         x_cv_final_counts = model.fit_transform(x_cv['CleanedText'].values)
         x_ts_final_counts = model.fit_transform(x_ts['CleanedText'].values)
```

```

# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

print("the type of count vectorizer ",type(x_tr_final_counts))
print("the shape of out text TFIDF vectorizer ",x_tr_final_counts.get_shape())
print("the number of unique words including both unigrams and bigrams "
, x_tr_final_counts.get_shape()[1])

print("the type of count vectorizer ",type(x_cv_final_counts))
print("the shape of out text TFIDF vectorizer ",x_cv_final_counts.get_shape())
print("the number of unique words including both unigrams and bigrams "
, x_cv_final_counts.get_shape()[1])

print("the type of count vectorizer ",type(x_ts_final_counts))
print("the shape of out text TFIDF vectorizer ",x_ts_final_counts.get_shape())
print("the number of unique words including both unigrams and bigrams "
, x_ts_final_counts.get_shape()[1])

```

```

the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (15000, 16311)
the number of unique words including both unigrams and bigrams 16311
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (5000, 9086)
the number of unique words including both unigrams and bigrams 9086
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (5000, 8627)
the number of unique words including both unigrams and bigrams 8627


```

```

In [60]: # dimensionality reduction for word2vec train data
c = sc.fit_transform(x_tr_final_counts)
truncatesvd_reduction(c)

```






Here 500 components explain 80% of the variation in word2vec train data

```
In [61]: svd = TruncatedSVD(n_components=500)
x_tr_dataset = svd.fit_transform(c)
x_tr_dataset.shape
```

```
Out[61]: (15000, 500)
```

```
In [62]: # dimensionality reduction for word2vec cross validated data
c = sc.fit_transform(x_cv_final_counts)
truncatesvd_reduction(c)
```




Here 500 components explain 80% of the variation in word2vec cross validated data

```
In [63]: svd = TruncatedSVD(n_components=500)
x_cv_dataset = svd.fit_transform(c)
x_cv_dataset.shape
```

```
Out[63]: (5000, 500)
```

```
In [64]: # dimensionality reduction for word2vec test data
c = sc.fit_transform(x_ts_final_counts)
truncatesvd_reduction(c)
```



Here 500 components explain 80% of the variation in word2vec test data

```
In [65]: svd = TruncatedSVD(n_components=500)
x_ts_dataset = svd.fit_transform(c)
x_ts_dataset.shape
```

```
Out[65]: (5000, 500)
```

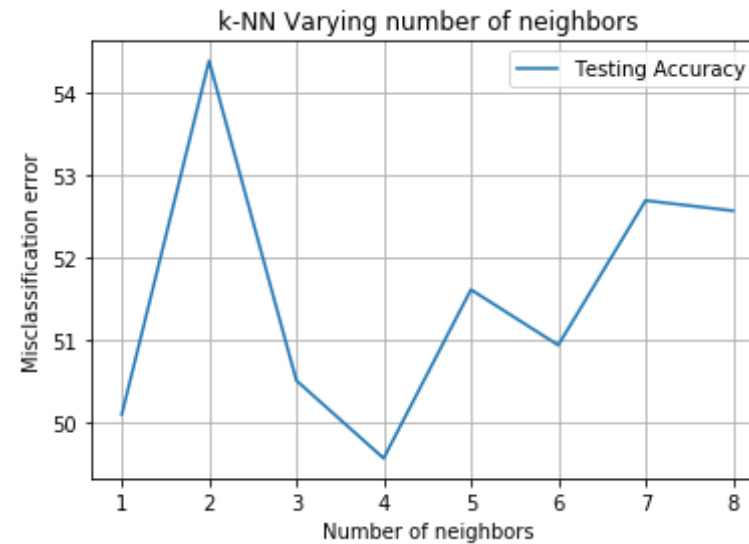
```
In [66]: print(x_tr_dataset.shape)
print(x_cv_dataset.shape)
print(x_ts_dataset.shape)
print(y_train.shape)
print(y_cv.shape)
print(y_ts.shape)
```

```
(15000, 500)
(5000, 500)
(5000, 500)
(15000,)
(5000,)
(5000,)
```

```
In [67]: # Finding Optimal K by simple Cross validation by kd- tree algo

optimal_k = find_optimal_k_kdtree(x_tr_dataset,y_train,x_cv_dataset,y_cv)
```

optimal\_k is:- 4



```
In [68]: from sklearn.metrics import f1_score
# ===== KNN with k = optimal_k =====
# instantiate learning model k = optimal_k

knn_optimal = KNeighborsClassifier(n_neighbors = optimal_k)

# fitting the model
knn_optimal.fit(x_tr_dataset, y_train)

# predict the response
pred = knn_optimal.predict(x_ts_dataset)


# evaluate f1_score
```

```
f1_w2c_kdtree = f1_score(y_ts, pred, average='macro') * float(100)
f1_w2c_kdtree_optimal_k = optimal_k
print('\nThe f1_score of the knn classifier of KDtree algo for k = %d is %f%%' % (optimal_k, f1_w2c_kdtree))
```

The f1\_score of the knn classifier of KDtree algo for k = 4 is 49.606452%

```
In [69]: skplt.plot_confusion_matrix(y_ts, pred)
```

```
Out[69]: <matplotlib.axes._subplots.AxesSubplot at 0x62a2f7b198>
```




```
In [70]: #classification report
print(classification_report(y_ts, pred))
```

|             | precision | recall | f1-score | support |
|-------------|-----------|--------|----------|---------|
| negative    | 0.11      | 0.20   | 0.15     | 554     |
| positive    | 0.89      | 0.81   | 0.85     | 4446    |
| avg / total | 0.80      | 0.74   | 0.77     | 5000    |

```
In [71]: # Finding Optimal K by simple Cross validation by brute algo

optimal_k = find_optimal_k_brute(x_tr_dataset ,y_train,x_cv_dataset,y_c
v)
```

optimal\_k is:- 4



```
In [72]: from sklearn.metrics import f1_score
# ===== KNN with k = optimal_k =====
# instantiate learning model k = optimal_k

knn_optimal = KNeighborsClassifier(n_neighbors = optimal_k)

# fitting the model
knn_optimal.fit(x_tr_dataset, y_train)

# predict the response
pred = knn_optimal.predict(x_ts_dataset)


# evaluate accuracy
f1_w2c_brute = f1_score(y_ts, pred, average='macro') * float(100)
```

```
f1_w2c_brute_optimal_k = optimal_k
print('\n\nThe accuracy of the knn classifier of brute algo for k = %d is
%f%%' % (optimal_k, f1_w2c_brute))
```

The accuracy of the knn classifier of brute algo for k = 4 is 49.606452%

```
In [74]: skplt.plot_confusion_matrix(y_ts ,pred)
```

```
Out[74]: <matplotlib.axes._subplots.AxesSubplot at 0x62bbbadfd0>
```



```
In [75]: #classification report
print(classification_report(y_ts, pred))
```

|             | precision | recall | f1-score | support |
|-------------|-----------|--------|----------|---------|
| negative    | 0.11      | 0.20   | 0.15     | 554     |
| positive    | 0.89      | 0.81   | 0.85     | 4446    |
| avg / total | 0.80      | 0.74   | 0.77     | 5000    |

```
In [79]: # TF-IDF weighted Word2Vec
#tfidf_feats = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(list_of_sent): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            # tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors.append(sent_vec)
    row += 1
```

100%|██| 15000/15000 [00:23<00:00, 62.6.98it/s]

```
In [80]: # TF-IDF weighted Word2Vec
#tfidf_feats = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectors1 = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(list_of_sent_cv): # for each review/sentence
```





```
100%|██████████| 5000/5000 [00:07<00:00, 70  
6.43it/s]
```

```
In [83]: # Finding Optimal K by simple Cross validation¶ by kd- tree algo

optimal_k = find_optimal_k_kdtree(x_tr_final_counts,y_train,x_cv_final_
counts,y_cv)
```

A line graph titled "k-NN Varying number of neighbors" showing the relationship between the number of neighbors (x-axis) and the misclassification error (y-axis). The x-axis ranges from 1 to 8, and the y-axis ranges from 49.5 to 52.0. The error starts at approximately 49.3 for 1 neighbor, rises to 50.4 for 2 neighbors, dips slightly to 50.3 for 3 neighbors, drops to 49.9 for 4 neighbors, peaks at 51.8 for 5 neighbors, dips to 51.4 for 6 neighbors, reaches its highest point at 52.1 for 7 neighbors, and ends at 51.6 for 8 neighbors.

| Number of neighbors | Misclassification error |
|---------------------|-------------------------|
| 1                   | 49.3                    |
| 2                   | 50.4                    |
| 3                   | 50.3                    |
| 4                   | 49.9                    |
| 5                   | 51.8                    |
| 6                   | 51.4                    |
| 7                   | 52.1                    |
| 8                   | 51.6                    |

```
In [84]: from sklearn.metrics import f1_score
# ===== KNN with k = optimal_k =====
# instantiate learning model k = optimal_k

knn_optimal = KNeighborsClassifier(n_neighbors = optimal_k)

# fitting the model
knn_optimal.fit(x_tr_final_counts, y_train)


# predict the response
pred = knn_optimal.predict(x_ts_final_counts)

# evaluate accuracy
f1_tfidf_w2c_kdtree = f1_score(y_ts, pred, average='macro') * float(100)
f1_tfidf_w2c_kdtree_optimal_k = optimal_k
print('\nThe accuracy of the knn classifier of KDtree algo for k = %d is %f%%' % (optimal_k, f1_tfidf_w2c_kdtree))
```

The accuracy of the knn classifier of KDtree algo for k = 1 is 48.942989%

```
In [86]: skplt.plot_confusion_matrix(y_ts ,pred)
```

```
Out[86]: <matplotlib.axes._subplots.AxesSubplot at 0x62c4c53390>
```



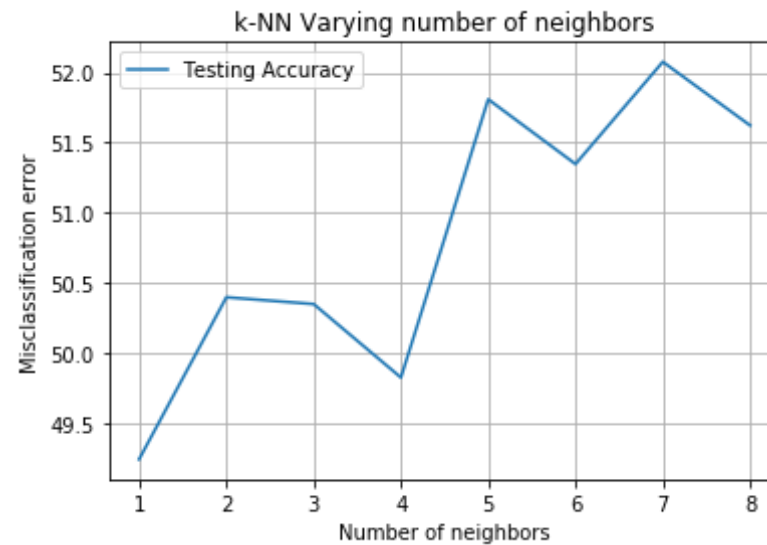
```
In [87]: #classification report
print(classification_report(y_ts, pred))
```

|             | precision | recall | f1-score | support |
|-------------|-----------|--------|----------|---------|
| negative    | 0.09      | 0.09   | 0.09     | 554     |
| positive    | 0.89      | 0.88   | 0.89     | 4446    |
| avg / total | 0.80      | 0.80   | 0.80     | 5000    |

```
In [89]: # Finding Optimal K by simple Cross validation by brute algo

optimal_k = find_optimal_k_brute(x_tr_final_counts,y_train,x_cv_final_c
ounts,y_cv)

optimal_k is:- 1
```



```
In [91]: from sklearn.metrics import f1_score
# ===== KNN with k = optimal_k =====
# instantiate learning model k = optimal_k

knn_optimal = KNeighborsClassifier(n_neighbors = optimal_k)

# fitting the model
knn_optimal.fit(x_tr_final_counts, y_train)


# predict the response
pred = knn_optimal.predict(x_ts_final_counts)

# evaluate accuracy
f1_tfidf_w2c_brute = f1_score(y_ts, pred, average='macro') * float(100)
f1_tfidf_w2c_brute_optimal_k = optimal_k
print('\nThe accuracy of the knn classifier of brute algo for k = %d is
%f%%' % (optimal_k, f1_tfidf_w2c_brute))
```

The accuracy of the knn classifier of brute algo for k = 1 is 48.942989%

```
In [92]: skplt.plot_confusion_matrix(y_ts ,pred)
```

```
Out[92]: <matplotlib.axes._subplots.AxesSubplot at 0x62dcf7a1d0>
```



```
In [93]: #classification report  
print(classification_report(y_ts, pred))
```

|             | precision | recall | f1-score | support |
|-------------|-----------|--------|----------|---------|
| negative    | 0.09      | 0.09   | 0.09     | 554     |
| positive    | 0.89      | 0.88   | 0.89     | 4446    |
| avg / total | 0.80      | 0.80   | 0.80     | 5000    |

```
In [100]: from prettytable import PrettyTable  
x = PrettyTable()  
x.field_names = ["Vectorizer", "Model", "hyperparameter(k)", "F1_Score"]  
x.add_row(["BOW", "KD-Tree", f1_bowt_kdtree_optimal_k, f1_bowt_kdtree])  
x.add_row(["BOW", "Brute", f1_bowt_brute_optimal_k, f1_bowt_brute])  
x.add_row(["TF-IDF", "KD-Tree", f1_tfidf_kdtree_optimal_k, f1_tfidf_kdtree  
)
```

```
x.add_row(["TF-IDF", "Brute", f1_tfidf_brute_optimal_k, f1_tfidf_brute])
x.add_row(["Word2Vec", "KD-Tree", f1_w2c_kdtree_optimal_k, f1_w2c_kdtree])
x.add_row(["Word2Vec", "Brute", f1_w2c_brute_optimal_k, f1_w2c_brute])
x.add_row(["TFIDF Word2Vec", "KD-Tree", f1_tfidf_w2c_kdtree_optimal_k, f1_tfidf_w2c_kdtree])
x.add_row(["TFIDF Word2Vec", "Brute", f1_tfidf_w2c_brute_optimal_k, f1_tfidf_w2c_brute])
print(x)
```

| Vectorizer     | Model   | hyperparameter(k) | F1_Score           |
|----------------|---------|-------------------|--------------------|
| BOW            | KD-Tree | 2                 | 50.761934          |
| BOW            | Brute   | 2                 | 50.761934          |
| TF-IDF         | KD-Tree | 1                 | 51.053415620344744 |
| TF-IDF         | Brute   | 1                 | 51.053415620344744 |
| Word2Vec       | KD-Tree | 4                 | 49.60645242557704  |
| Word2Vec       | Brute   | 4                 | 49.60645242557704  |
| TFIDF Word2Vec | KD-Tree | 1                 | 48.942988955975764 |
| TFIDF Word2Vec | Brute   | 1                 | 48.942988955975764 |

As per the above chart observation, "word2vec KD-tree" and "word2vec Brute" is well fit. others Bow , TF-IDF and TFIDF word2vec models are underfit. so we will consider "Word2vec KD-tree" and "Word2vec brute" model for classification positive and negative datapoints with hyperparameter 4.